Seven research questions, seven evidence contracts

A map of local evidence and its limits | No common effect scale or overall ranking

	Phenomenon	Local evidence	Main finding	Evidence status	Critical qualification
RQ1	Ranking utility	G1 primary LLM ranking; local traditional comparison	Utility varies by dataset, model and prompt. HR@5/NDCG@5 admit local numerical comparison.	Adapted; locally comparable	Shared positive-forced 20-item pool; no universal winner. Traditional-vs-LLM MRR depths differ (20 vs 10). Not paper-exact.
RQ2	Beyond-utility dimensions	Corrected G1 metrics; traditional counterparts where defined	Long-tail, popularity, novelty, diversity and coverage vary independently by configuration.	Adapted + local extensions	Traditional: full cohort; G1: strict-valid/defined users. Calibration is a local extension; Food.com PopREO is undefined.
RQ3	History-length sensitivity	Completed GE20 trajectory; separate paired G2 contrasts	Early GE10 H10−H0: 57 pos / 1 neg / 32 inc. Later GE20 H20−H10: 3 pos / 18 neg / 69 inc.	Controlled extension	Contrast counts use different eligible cohorts. H0 is controlled no-history on warm users; no monotonic benefit or universal optimum.
RQ4	Candidate-position sensitivity	G3 paired FIRST vs RANDOM order stress test	FIRST−RANDOM: 85 pos / 1 neg / 4 inc across 90 configuration×metric contrasts.	Controlled stress test	FIRST places the known target first: target-informed oracle condition, not a deployable ordering policy.
RQ5	Generated user profiles	Corrected matched History-only / Profile-only / Profile+History; systematic support	Profile effects are configuration-dependent; some gains occur, but no general benefit is established.	Matched local evidence	No validated semantic-fidelity claim or universal causal profile-benefit claim.
RQ6	Output reliability	G1–G4 constrained candidate-index structural-validity accounting	Strict structural validity spans 98.5–99.1% across G1–G4.	Adapted construct	Paper free-text textual/catalogue hallucination is not local structural invalidity; direct reproduction is absent.
RQ7	Two-stage reranking	Canonical F1BZQ candidate pools + G4 conditional/unconditional utility	Candidate Recall@20 spans 0.4–46.9%. End-to-end HR@5/NDCG@5 factor into recall × conditional utility.	Pipeline decomposition	Single-positive identity only. No same-pool superiority test; Food.com conditional n=4 is diagnostic.

Statuses describe study design and evidence scope; they are not quality grades.